

The Riscy Processor

Karan Bavishi, Mark Mansi, Suhas Pai, Preyas Shah

1 Introduction

We implement the RV64I set of instructions in the RISC-V ISA[1] with a superscalar, out-of-order core. Our implementation is optimized for correctness. Our second goal was performance. Our third goal was synthesizability.

In Section 2, we provide a brief overview of our design. In Section 3, we talk about the design of each individual module in detail. Section 4 describes some of the policies we designed and implemented to boost the performance of our processor. In Section 5, we discuss and explain some interesting design decisions we made. In Section 6, we talk about a few challenges that we ran into, which forced us to rethink portions of our design. Finally, in Section 7, we discuss the performance achieved by our processor for certain benchmarks, along with some key take-away points we learned by our experience.

2 Overview

- 4-wide pipeline
- Out-of-order execution; In-order commit
- Fetch (3 cycles for cache hit)
 - Branch prediction (TODO): 2-level predictor, BTB, global history register.
 - I-Cache: Blocking, 4 KB, 2-way associative, 1 read/write port, 2-cycle latency.
 - Prefetch Buffer: 4 entries, 32 bytes each
- Decode and Rename (2 cycles)
 - 64-entry ROB
- Issue and FUs (3 cycles)
 - 4 ALUs
 - 4 issue queues, 16 entries each, 1 ALU per queue
 - Load-balancing arbiter places new instructions in issue queues.
 - ALUs also used to compute `ld/st`/branch addresses.
- Address Queue
 - 32 entries
 - Tracks load/store dependencies.
 - Non-speculative memory disambiguation.
 - Loads access cache out-of-order between stores.
 - Stores access cache on commit.
 - D-Cache (TODO)
- Writeback (2 cycle)

- Processor supports back-to-back execution of dependent instructions.
 - Writeback structure (a.k.a ROB WB or more affectionately, FooPP) is designed to avoid the massive tangle of wires created by broadcast-based writeback among 4 ALUs and a LSQ.
- Commit (1 cycle)
 - Stall: we implement a distinct top-level module to handle stall coordination among all stages. The two stall producers in the pipeline are the issue queue arbiter and the ROB. A stall is generated when there is not enough room in the issue queues or the ROB. Each stage is a consumer of stalls produced by stages before it.

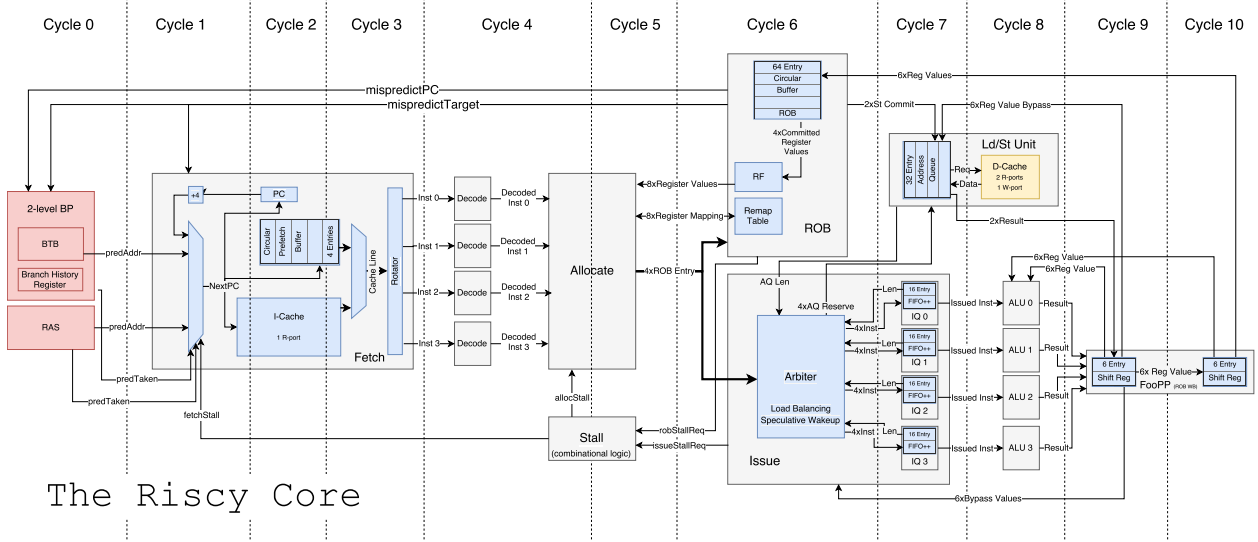


Figure 1: *Top-level diagram of our processor* A larger version of the diagram is also available at [2].

3 Modules

Our processor has an 11 cycle pipeline depth, as shown in Figure 1. The pipeline is divided into several independently designed and implemented blocks (implemented as modules in Chisel). This section presents an overview of each block.

3.1 Fetch

Our fetch module encapsulates the logic of selecting the next PC, issuing requests to the instruction memory hierarchy (including the I-cache), and presenting raw instructions to the subsequent stage (decode).

We implement a pipelined, blocking, 2-way associative instruction cache with 2-cycle latency. An additional cycle of latency comes from computing the next PC.

The instruction cache returns a hit if it finds the requested entry. If not, it issues a refill request to memory for the missing address. Once the memory responds, it refills its data and tag arrays with the response. The cache line replacement policy is random replacement.

We also added a prefetch buffer, containing 4 entries. It does cache block prefetching; i.e. it stores the additional cache lines sent by memory as a part of its response to a refill request. This is primarily done to get around the single-port limitation of the I-cache. This prefetch buffer can be easily changed to do next-line prefetching as a part of future work.

The output from the instruction cache is rotated and presented to the rest of the pipeline.

3.2 Decode and Allocate

The next two cycles decode and rename the instructions from fetch.

Decode takes less than one cycle because it contains no logic (it just separates wires). Thus, we combine decode with the first cycle of renaming.

Renaming proceeds in two phases. The first renames all destinations. This can be done regardless of which instruction is being renamed because all instructions get a single ROB entry. The second cycle renames operands and produces an ROB entry which can just be latched to the ROB next cycle. Renaming operands requires the knowledge of the renamed destinations and previous remappings from the remap table. Producing an ROB entry, requires reading the register file and ROB because our design is read-before-issue.

3.3 ROB

The ROB is implemented as a 64-entry circular buffer. We choose to use an ROB as opposed to a merged register file because we prioritized performance and ease of implementation over energy. It serves a few major roles in the pipeline.

1. It serves as a central location of information about instructions currently in the pipeline, including operands, renaming information, and written-back values. Thus, the ROB must be read by Allocate stages to produce new ROB entries. In any given cycle, up to 4 new instructions may be entered into the ROB.
2. It is responsible for updating the architectural state when an instruction commits, including updating registers and notifying the L/S Unit that a store should go to memory. Up to 4 instructions can commit in a cycle, including up to 2 store instructions.
3. It serves as a point of serialization for instructions to commit in order. This preserves interruptability and allows recovery from mispeculations.
4. It notifies the rest of the pipeline when a mispeculation happens and is responsible for communicating to other modules which instructions in the pipeline are valid. It also notifies fetch and the branch predictor of the correct branch target.

Eras To aid with identifying squashed instructions, the ROB keeps track of the mispeculation *era*. New instructions coming into the pipeline are deemed to be part of the current era. When a mispeculation occurs, a new era begins. Instructions in the pipeline are uniquely identified by their ROB entry number and tag. These values are propagated through the pipeline with the instruction. The ROB also provides the current era number to other pipeline stages.

3.4 Issue and FUs

In the same cycle, an ROB entry both latches in the ROB and goes to the issue queue arbiter, which puts it in a queue. All instructions go through the arbiter regardless of their type. Notably, the issue queues are designed to be largely oblivious to a number of factors which would have complicated implementation (and the issue queues are already highly complicated structures). For example, the issue queues are mostly unaware of the current ROB era and even of instructions' opcodes.

Our issue stage consists of an arbiter feeding 4 issue queues. The arbiter does load balancing to attempt to keep any queue from becoming filled while other ALUs are available.

Each issue queue is independent and feeds a single ALU. The issue queues can each independently wake up a single instruction to be sent an ALU. Our implementation is capable of speculatively waking up ALU operations when their operands are known to be nearly ready. This allows our processor to execute instruction with back-to-back dependencies. Our writeback mechanism helps with this also...

3.5 Writeback

Each cycle, our processor may produce up to 6 values that need to be written back (4 from ALUs and 2 from memory). To avoid creating $O(n^2)$ wires between all of these value producers/consumers, we designed a central writeback structure, which we call **FooPP**. All values are written to the **FooPP** in the cycle they produced. Any consumers of those values, including the ROB, issue queues, and L/S Unit, take the values directly from the **FooPP**.

To avoid adding a 1 cycle bubble between instructions with back-to-back dependencies, the **FooPP** keeps 1 cycle of history, which the FUs can then consume from easily.

The writeback structure also filters out any values coming from squashed instructions by comparing the era of the input values with the current ROB era number. This allows us to get rid of the squashed instructions easily by letting them percolate through the pipeline.

3.6 Load/Store Unit

When a load or store instruction reaches the arbiter, the arbiter asks the address queue to reserve a spot for it. If the address queue is full, the pipeline stalls. The arbiter then puts the load or store in a normal issue queue to wait until an ALU can compute its address. When the address is written to the **FooPP**, the address queue will consume it.

Loads can access the D-cache as soon as their address is known and it is also known that there is no preceding store to the same address. Stores can access the D-cache as soon as the ROB signals the address queue that they have committed.

To limit the number of ports need on the D-cache, we assume only two ports. The ROB gives the guarantee that no more than 2 stores will commit in any cycle.

4 Performance Enhancement Policies

In this section, we talk about some of the design decisions we took which were aimed at improving the performance of our processor.

4.1 Fetch

We added a prefetch buffer which stored cache blocks when a refill response was received from memory. This helped us get around the limitations of having a single-ported I-cache, which allowed us to only write one cache line when the refill response was received. Having a prefetch buffer allowed us to store multiple cache lines.

The prefetch buffer also helped improve performance by having enough instructions to keep the backend busy. Without the prefetch buffer, we were only able to serve instructions every 2 cycles before incurring another hit.

4.2 Decode and Allocate

Decode is combined with the first stage of Allocate to avoid wasting a cycle.

4.3 ROB

In case of a misprediction, there are two ways to squash the speculated instructions. The first involves adding valid bit wires everywhere, and resetting them when a misprediction occurs. The second way is to let the squashed instructions percolate through the pipeline normally and ignore the result. We chose the second approach in our design.

However, the ROB cannot latch new instruction unless all the squashed instructions have been flushed from the system because ROB entry numbers would cease to be unique (a correct instruction and a squashed instruction may alias). Thus, we add misprediction eras. As stated previously, Whenever ROB realizes a misprediction, it increments the era counter, so squashed instructions are those from previous eras.

This allows our pipeline to begin execution of correct instructions before the pipeline is fully flushed.

4.4 Issue

Renamed instructions are sent to both the ROB and Issue modules in parallel to save a pipeline stage.

We also added support for speculative wakeup of instructions whose operands would be ready in the subsequent cycle. This allowed us to perform back-to-back execution for dependent instructions and boost performance. (At the moment, we only support speculative wakeup for ALU operations, but this logic can be extended for memory loads, too).

4.5 Execute

The FooPP module is also responsible for storing the results for 2 cycles for bypass. This bypass support allowed us to execute dependent instructions in back-to-back cycles.

5 Design Decisions

In this section, we talk about the interesting design choices we made for our processor.

1. Pipelined Decode & Rename - In our Decode and Allocate design, we decode and rename the destination in the first cycle. Renaming of operands is done in the second cycle. This allows us to allocate new entries in the ROB in the second cycle, in parallel with the renaming of operands.
2. Use Writeback to flush values from squashed instructions - As mentioned in Section 4, we decide to let the instructions to be squashed percolate through the pipeline. An interesting question was to how to filter out the results from these squashed instructions and not make them be mistakenly used for bypass. We ultimately decided to add some filter logic to the Writeback module inputs. This made sure that the squashed values were invalidated and thus never available for bypass.
3. Arbiter load-balancing and bypass - Our original idea was to have the arbiter build dependency graphs and put dependent instructions in the same queue. The rationale was that putting dependent instructions in the same queue would reduce the overhead of bypass across queues. However we soon realized that this plan would result in a multi-cycle arbiter design.

We ultimately decided to have a simple arbiter which would issue instructions based on load balancing, i.e. instructions would be issued based on the emptiness of the queues. This allowed us to design a single-cycle arbiter. However, this brought back the original problem of complex bypass. For this, we designed the Writeback structure which would allow bypass of computed values to the 4 ALUs.

4. LSQ: Use Issue queue wakeup logic for load/store instructions - Our original plan was to have 4 issue queues and 1 load-store queue, using 4 ALUs and 1 AGU. But we decided to re-use the existing ALUs for address computation for loads and stores. This would allow us to re-use the complex instruction wakeup logic already present in the Issue Queues, and not re-implement from scratch.
5. LSQ: Serialize loads and stores so that D-cache can have 1 read port and 2 write ports - TODO
6. Unified Stall logic - We decided to come up with a unified stall module which would take stall requests as input from ROB and Arbiter modules and generate stall outputs to the consumers such as the Fetch and the I-cache modules. This unified logic simplified the interface between all these modules, and avoided complex wiring.

6 Design Challenges

In this section, we talk about certain challenges that we encountered which forced us to modify our design.

6.1 Cancelling memory requests

After adding misspeculation recovery support, we found that our processor had a 5-cycle bubble before instructions were fetched from the new target. On investigating more, we realized that it was because the I-cache was busy requesting an address from memory that would no longer be needed. To get around this, we added support for cancellation of requests to memory. This required changes to the memory and I-cache interfaces.

6.2 ROB deadlock issues

During testing, we found the processor often entered a state of deadlock. We realized that we needed to rethink the mean of stalling the ROB. We found that the ROB should conceptually be thought of as two independent modules: allocation and commit. Allocation should stall. But since commit is the only way for the processor to make progress, it should never stall.

6.3 Issue queue timing

Each issue queue is FIFO-like. Because of writeback and wakeup logic, entries in the middle of the queue may be woken up and removed or writtenback to and modified. In addition, an entry may shift while it is being modified. Finally, the issue arbiter is inserting new entries into the issue queues at the tail.

Because of the high fan-in, we had to carefully design the issue queue logic to ensure robustness to timing issues. If not done carefully, entries could be duplicated or lost, fail to be issued or fail to be updated.

6.4 Zombies

While adding support for misspeculation recovery through percolation, we realized that squashed “zombie” instructions could end up staying much longer in the issue queue if their operands not ready. Consider the case where a zombie load instruction is kept waiting for 100 cycles waiting for the memory to respond to the request. This can happen because in our percolation policy, we execute the zombie instructions but decide to filter the results, simplifying issue logic.

In order to mitigate this problem, we allow instructions to issue immediately if they are zombies, even if their operands are not ready. Since we do not care about the results of these instructions, we allow them to wakeup as soon as an ALU is available.

7 Performance

7.1 Methodology

We ran a series of workloads to show the upper and lower bounds of IPC that our processor could achieve. We ran the workloads in two main settings to test the impact of lack of a data-cache and data-prefetching:

- 7 cycle memory latency
- 4 cycle memory latency (smallest possible latency supported by our design)

Additionally, we also ran the same workloads with instruction prefetching turned on. In Section 7.3, we talk about the results achieved with the two above-mentioned settings, and along with instruction-prefetching wherever correctness of the benchmarks was verified.

7.2 Benchmarks

These are the benchmarks we ran:

Benchmark	ILP	Instructions	Description
Counters (optimal)	High ILP	8K insts	32 counters incremented independently.
Fibonacci	Max ILP 1.5	150 insts	Compute the first 64 Fibonacci numbers.
Daxpy	High ILP (50 % memory insts)	3K insts	Element-wise addition of arrays.
Stall (worst case)	Low ILP	160 insts	Back-to-back dependent instructions force in-order execution

Briefly, here is why we chose each benchmark:

1. *Counters* - In this benchmark, we increment 32 counters independently. This benchmark has extremely high ILP, so it is optimal for a superscalar out-of-order processor. This benchmark reveals the upper limit of IPC that can be achieved by our design.
2. *Fibonacci* - The Fibonacci benchmark is aimed at testing our design with a purely compute-based workload. Due to the dependencies in the Fibonacci recurrence relation, the maximum ILP that can be achieved is 1.5 IPC.
3. *Daxpy* - We use a loop-unrolled version of the DAXPY benchmark. Since the original version involves multiply instructions which are not supported in our design, we simply add instead. This benchmark has a high ILP workload with a high memory-to-compute ratio (50%) of instructions. This benchmark shows how much the lack of data-caching affects our performance.
4. *Stall* - In this benchmark, we execute 160 back-to-back dependent instructions. The maximum ILP that can be achieved is 1. This benchmark verifies that our design can execute dependent instructions in back-to-back cycles with no bubbles.

7.3 Results

In the following table, we describe the IPC achieved for all the benchmarks in the 2 chosen memory latency settings, with instruction prefetching enabled and disabled. We omit results for settings where we were unable to verify correctness.

Setting	Counters	Fibonacci	Daxpy	Stall
7 cycle, I-prefetch disabled	0.80	0.74	0.80	0.74
7-cycle, I-prefetch enabled	2.22	1.30	-	0.91
4 cycle, I-prefetch disabled	1.14	1.04	0.88	0.92
4 cycle, I-prefetch enabled	2.66	-	-	-

The following interesting observations can be made:

1. *Instruction prefetching makes a huge difference.* As shown by the IPC of the best-case Counters benchmark, it is hard to achieve a high throughput unless you have enough instructions to keep the backend busy. With instruction prefetching disabled, our processor incurs the memory latency penalty every 2 cycles.

Interestingly, even with prefetching enabled, our design is only able to reach an ILP of 2.66. The reason for this is that our prefetching design is fairly naive; it is not even a next-line prefetcher. The prefetch unit does not fetch instructions well in advance and only stores additional cache lines when a penalty is incurred. If we increase the size of the prefetch buffer from 4 to 8, our projections show that we can achieve an IPC of about 3.2.

Having a more intelligent prefetcher and increasing the size of the prefetch buffer should help boost the IPC close to the maximum theoretical limit of 4.

2. *It is important to have a big instruction window.* Similar to our previous observation, it is extremely important to have a big instruction window to exploit the techniques implemented in an out-of-order design. Without enough instructions in the frontend, it is difficult to find enough ILP, and the processor remains hampered by dependencies and memory latency.
3. *Memory latency can significantly degrade the IPC.* As can be seen from the IPC numbers achieved by the Daxpy benchmark, it is hard to achieve a high ILP unless the memory latency can be hidden effectively.

The Daxpy benchmark has an inherently high-ILP workload, but since our design lacks a data-cache, it is unable to go beyond executing more than one instruction per cycle.
4. *Close to optimal performance in Fibonacci* - We achieve an ILP of 1.30 for the 7-cycle, prefetching enabled setting, which is quite close to the theoretical limit of 1.5. We can actually reach an even higher figure of 1.42 for the 4-cycle, prefetching enabled setting, but we decided not to show it here because we couldn't get the correct result from the benchmark.
5. *We have near optimal performance in Stall.* Our Stall benchmark results are very close to the optimal IPC of 1. This highlights our processor's ability to execute dependent instructions in back-to-back cycles with no bubbles.

References

- [1] RISC-V ISA - <https://riscv.org/specifications/>.
- [2] RISCY Core diagram - https://github.com/mark-i-m/riscy/raw/master/riscy/doc/riscy_diagram.pdf